

AI-POWERED PROFIT INTELLIGENCE

Helping marketplace sellers understand their business through AI

Built from 0→1 to help Indian marketplace sellers replace spreadsheets with AI-powered profit insights.

OWNED END-TO-END

Research

Product Strategy

UX

AI Design

Roadmap

Launch

Iteration

Also built the initial MVP and shipped UI fixes & feature / bug resolutions.

ROLE


Product Manager • 0→1

AI STACK

GPT-4o-mini • Prompt Engineering • Grounded Context • Voice AI • Python

TIMELINE

4-Month Build • 2026



30+

Active Sellers

₹60L+

GMV Tracked

2L+

Orders Analysed

Amazon India + Meesho • early estimate, not yet verified

How SellerOps AI went from idea to production

DISCOVER

The Opportunity

The problem, the user, and why it was worth solving

UNDERSTAND

Understanding the User

Research, personas, and the seller journey

DECIDE

Strategy & Prioritization

Goals, key decisions, and the tradeoffs behind them

BUILD

Building trustworthy AI

Architecture, grounding, evaluation, guardrails & human review

VALIDATE

Shipping the Product

Roadmap, the shipped product, and a live demo

GROW

Impact & Learnings

Honest metrics, impact, and what comes next

SECTION 01 / DISCOVERY

01

Understanding Why Sellers Don't Trust Their Numbers

Every business decision depended on spreadsheets that couldn't answer one simple question: "Am I actually profitable?"

A workspace for true profit - not another dashboard

Sellers upload their marketplace data exports and get a clear picture of **true profit** - after fees, returns, shipping and cost of goods - alongside an AI assistant that explains what’s happening and what to do next.

Instead of wrestling with settlement reports and spreadsheets, an owner-operator can ask a question in plain language and get an answer grounded in their actual numbers, in English or Hinglish.

- AI Copilot
- Profit Intelligence
- Reconciliation
- COGS Tracking

EARLY-STAGE SCALE

Early Users	30+
GMV Analysed	₹60L+
Orders processed	2L+

Early estimates from a pre-launch / early-access stage — directional, not yet verified.

“After all the fees, ads and returns...**am I actually making money, and on which products?**”



01

PRIMARY PAIN

Reconcile the money

Orders vs. settlements vs. refunds – each in a separate export.

Orders · Settlements · Refunds



02

Decode marketplace fee complexity

FBA, referral, TCS/TDS, advertising, ABA and more.

6+ fee types stacked



03

Identify loss-making SKUs

Loss-makers hiding behind healthy top-line revenue.

Hidden in top-line



04

Do it every week

Hours in Excel, weekly – and still missing deductions.

~5-15 hrs / month

TODAY CSV exports → Excel → Manual calculation → Uncertain decisions



WITH SELLEROPS AI → Trusted profit → Confident decisions

Sellers can't turn analytics into trustworthy action

CURRENT SOLUTION	WHY IT FALLS SHORT
Native Marketplace Reports	Raw data, dozens of columns, no unified profit view
Excel / Sheets	High effort, error-prone, hard to maintain weekly
Traditional BI Platforms	Require setup; don't understand marketplace fee semantics
Accounting software	Doesn't understand marketplace fees, returns, or SKU-level profitability
Generic ChatGPT 	Hallucinates numbers that don't match actual settlements

CORE PROBLEM

Existing tools either overwhelm sellers with data or generate answers they can't trust.

Without trusted profit data, sellers scale the wrong products, cut the wrong costs, and make poor inventory decisions.



PRODUCT INSIGHT

Sellers don't need more analytics. They need **trusted answers grounded in their own business data.**

02

Understanding the User

Instead of relying on assumptions, we used demos, support conversations and early user feedback to shape every major product decision.

Customer Interviews

Research Insights

Personas

Journey Mapping

Research That Changed the Roadmap

Early-stage research · Customer discovery · Pre-launch insights · Discovery phase

INSIGHT	EVIDENCE	DECISION
 Sellers trust settlements, not order revenue	P&L disputes surfaced in demos	Ground AI in settlement logic; default to “last complete month”
 Users won’t read long reports	Low engagement on static tabs	Prioritized AI Chat as the primary product experience
 Blank-page anxiety stalls first use	Hesitation on an empty chat box	Added guided prompts to improve first-session activation
 Off-topic queries waste quota	“Rewrite my listing” requests	Scope guardrails & forbidden-topic filters
 Meesho economics ≠ Amazon	RTO/COD dynamics differ structurally	Separate Meesho system prompt & context builder
 KEY TAKEAWAY Every major feature in SellerOps AI traces back to a validated customer insight - not an assumption.		

“

I see ₹2L in Seller Central but my bank account says something else. **Just tell me what's wrong.**

Yuvraj Arora, Amazon Seller (paraphrased from customer interview)

Three sellers we build for

R

Rahul · 32

Amazon FBA Founder · Jaipur

GMV

₹5L/mo

MARKET

Amazon

TECH

Medium

Men's apparel brand · manages own ads · checks Seller Central weekly

"I know sales are growing. I don't know whether profits are."

GOAL

Grow profit without spending weekends in Excel.

PAIN

Doesn't trust settlement reports.

JOB TO BE DONE

Tell me which SKUs deserve more inventory.

P

Priya · 29

Meesho Ops Manager · Surat

GMV

₹8L/mo

MARKET

Meesho

TECH

Low-Med

Fashion & sarees · high COD volume · runs daily label & return workflows

"Every return costs money and nobody tells me why."

GOAL

Keep operations smooth while protecting margin.

PAIN

Returns & RTO silently eat her profit.

JOB TO BE DONE

Help me reduce COD & return losses.

A

Amit · 38

Multi-Marketplace Seller · Delhi

GMV

₹20L/mo

MARKET

Amazon, Meesho, Flipkart

TECH

High

Home & kitchen multi-brand · sells on both platforms · outgrew spreadsheets

"I outgrew spreadsheets two years ago."

GOAL

Run every channel from one clean cockpit.

PAIN

No single view across marketplaces.

JOB TO BE DONE


Optimize multiple marketplaces from one place.

SellerOps AI

02 · Understanding the User

Customer Journey Map

Goal · Emotion · Pain · Opportunity · Metric



OUR BIGGEST OPPORTUNITY The largest drop-off was **report upload during setup** - so it became the **highest-priority product investment** on the roadmap.

03

Turning Insights into Strategy

Customer research revealed dozens of opportunities. We focused on the few that built trust fastest.



North Star



Trade-offs



Prioritization Framework

"Great products are defined as much by what they don't build as by what they do."

★ NORTH STAR METRIC

Weekly active sellers who receive an actionable AI insight and take a follow-up action.

Why actions, not conversations: conversations don't create business value — decisions do.

AI GUARDRAILS

Hallucination rate	< 2%
Grounded answer rate	> 98%
Median response	< 8 s
Cost / insight	< \$0.08

LEADING INDICATORS

ACTIVATION

≥40%

ask ≥1 question in first 14 days

ENGAGEMENT

≥3

AI questions / week per user

VALUE

≥25%

suggested actions clicked / session

MONETIZATION

≥15%

quota-exhausted trials → upgrade



BUSINESS OUTCOMES

Retention

Revenue

Trust

Insight → Trust → Action → Retention

Our Product Strategy

PRODUCT THESIS

If we ingest a seller’s existing reports and compute settlement-aligned true profit, then wrap it in conversational AI, **sellers will repeatedly use the product, trust the insights, and ultimately pay for ongoing clarity – not just one-time reconciliation.**

OUTCOME This strategy deliberately prioritized **trust over feature breadth** – producing an MVP focused on profitable decision-making, not comprehensive analytics.

MVP PHILOSOPHY

-  **Accuracy before polish**
Fast time-to-value over deep integrations
-  **Time-to-first-insight over integration depth**
Upload → dashboard, not connect API → wait → configure.
-  **Earn trust before expanding scope**
Meesho, label cropper and ads analytics came after the profit core was credible.

⌵ DELIBERATELY DEFERRED

- Live marketplace APIs
- Multi-user collaboration
- Custom dashboards
- Ads optimization
- Inventory forecasting

CUSTOMER PROBLEM	PRODUCT RESPONSE
“Am I profitable?”	True profit after all fees & COGS
“Which to kill?”	SKU margin, return loss, alerts
“Why did profit drop?”	AI analysis + recommendations
“Where’s money leaking?”	Reconciliation, RTO, fee breakdown

Four decisions, four tradeoffs made on purpose

D1 Report upload

chosen over **live marketplace APIs**

● WHY IT WON

Ships fast, no OAuth — and Meesho has no accessible API.

● TRADEOFF

Manual uploads, no real-time sync, onboarding friction.

IMPACT Amazon + Meesho supported at launch, zero integration approvals.

D2 AI-first chat

chosen over **static dashboards**

● WHY IT WON

Differentiates from spreadsheets; easy for non-technical sellers.

● TRADEOFF

Inference cost, hallucination risk, LLM dependency.

IMPACT AI chat is the primary reason sellers return week over week.

D3 Trust before growth

accuracy chosen over **feature breadth**

● WHY IT WON

One wrong number destroys trust in a finance-adjacent product.

● TRADEOFF

Slower velocity, fewer features, heavier backend.

IMPACT Profit numbers reconcile to settlements – the product's trust anchor.

D4 Grounded AI

chosen over **generic prompting**

● WHY IT WON

Answers built from the seller's own data stay factual, not invented.

● TRADEOFF

Context-pipeline complexity, higher latency per query.

IMPACT Near-zero fabricated numbers – answers match the dashboard.

Six ways to deliver AI insight - one won

EVALUATED ON Accuracy · Cost · Time-to-value · Eng effort · User trust

OPTION A

Static weekly AI reports

× Rejected

PROS

Cheap · easy to build · simple onboarding

CONS

Static · low engagement · no follow-ups

OPTION D

Rule-based chatbot

× Rejected

PROS

Predictable · cheap · reliable

CONS

Can't answer new questions · doesn't scale

OPTION F

Voice-first mobile app

→ Deferred

WHY NOT NOW

No demand signal in interviews · desktop workflows first · low MVP ROI.

OPTION B

RAG over raw CSVs

× Rejected

PROS

Flexible · minimal preprocessing

CONS

Hallucinates math · costly context · weak on big tables

OPTION E

Copilot on every tab

× Rejected

PROS

Context-aware · familiar UI

CONS

Fragmented · repeated prompts · costly to maintain

OPTION C

Grounded chat

✓ Selected for MVP

Pre-aggregate the business → inject as context → LLM interprets. Numbers stay deterministic; language stays natural.

WHY IT WON

Deterministic calcs

Grounded answers

Natural language

Easy to extend

TRADEOFFS

More backend work · context pipeline · ongoing prompt engineering.

How we prioritized: $\text{Impact} \times \text{Trust} \div \text{Cost}$

- I

Impact

Does this help a seller make a better business decision?
- T

Trust

Can users rely on the answer without re-verifying it?
- C

Cost

Can we build and serve this sustainably?

DECISION PRINCIPLE

When Impact and Trust conflicted, Trust always won.

INITIATIVE	IMPACT	TRUST	COST	PRIORITY
Grounded AI chat				P0
Meesho-specific prompt				P1
Listing Generator AI				P2
Autonomous actions				Deferred
Why deferred: sellers weren't yet confident enough in AI recommendations to let the system act automatically - a direct read on the Trust axis.				

04

The AI System

Accuracy mattered more than model sophistication, so we grounded every answer in verified seller data before asking the LLM to reason.

Architecture

AI Evaluation

Guardrails

Grounded data

→

Context

→

LLM

→

Verified answer

We never asked the LLM to **calculate profit**. We asked it to **explain verified calculations**.

Grounded context, not retrieval

Because every answer depends on a **small, trusted financial context** – not millions of documents – we optimized for **determinism over retrieval**.

WHY TRADITIONAL RAG WASN'T THE FIT

- ✗ **No vector store**
Knowledge doesn't change between questions.
- ✗ **No embeddings**
Business metrics are structured, not semantic.
- ✗ **No retrieval at query time**
The whole context already fits in one prompt.
- ✗ **No LLM-generated SQL**
Avoid non-deterministic financial calculations.

Tradeoff: can't answer arbitrary queries outside the pre-computed context – custom date ranges or long-tail SKUs.

THE PIPELINE

Seller uploads reports



Python reconciles numbers



Structured business context



GPT-4o-mini explains



Actionable answer

WHY IT MATTERED


- ✓ Same numbers everywhere — chat matches dashboard
- ✓ No hallucinated calculations
- ✓ Fast responses, lower inference cost
- ✓ Easier to evaluate & no SQL-injection risk

THE KEY DETAIL

The LLM never performs arithmetic. Every financial calculation happens **before** inference.


From raw report to reviewed decision

All calculations happen **before inference** – the model only interprets trusted business context. Highlighted stages are the core value chain.



Deterministic calculations → Grounded context → Natural-language reasoning → Human review

What goes in, what comes back



SellerOps AI vs. retrieval-first AI

CAPABILITY	SELLEROPS AI	RETRIEVAL-FIRST AI
Uses seller business data	✓	✓
Runtime vector retrieval	×	✓
LLM calculates financial metrics	×	Often
Profit computed before inference	✓	×
Grounded business context	✓	Partial
Deterministic business rules	✓	Partial
Human validates decisions	✓	Optional

WHY THIS ARCHITECTURE In finance-adjacent products, **predictable numbers matter more than flexible retrieval**. SellerOps trades generality for trust.

Evaluation: current & next

Consciously sequenced, not skipped

WHY LATER As a 0→1 product, we prioritized shipping **trustworthy business logic** before investing in automated AI-evaluation infrastructure.

EVALUATION MATURITY	
Golden Q&A dataset	Planned
Automated regression tests	Planned
Offline benchmark suite	Planned
Human evaluation loop	Manual today
Usage & quota metering	Live

HOW WE VALIDATE TODAY


✓ Outputs spot-checked against settlement reports

✓ Prompt changes validated on historical seller cases

✓ Manual review vs. dashboards before release

✓ Anti-hallucination hints baked into context

If settlement data is missing, the assistant says “Insufficient data to calculate true profit” instead of estimating.



⚠ CURRENT RISK Prompt changes need manual validation — no automated regression guards against silent answer drift.

Failure modes - cause, response, owner

⚠️ Hallucinated numbers Cause: generic models invent figures. Mitigation: answer only from settlement summaries. OWNER · DATA CRITICAL	⚠️ Wrong COGS Cause: missing / estimated cost inputs. Mitigation: warn user + request missing COGS. OWNER · USER INPUT CRITICAL	Missing / stale uploads Cause: no recent reports; data drifts. Mitigation: prompt to upload latest report. OWNER · UX MEDIUM
Incomplete context Cause: SKU outside the pre-aggregated set. Mitigation: state the limitation, don't guess. OWNER · DATA MEDIUM	Ambiguous questions Cause: no clear date range or marketplace. Mitigation: ask a clarifying follow-up. OWNER · PROMPT LOW	Out-of-scope queries Cause: e.g. "rewrite my listing." Mitigation: politely decline & redirect. OWNER · PROMPT LOW

KEY INSIGHT

Most expected failures originate from **missing, stale or incomplete business data** – not from the language model itself.

WHEN CONFIDENCE IS LOW


Ask a follow-up · explain missing inputs · refuse to estimate.

The AI proposes; the seller decides

GUARDRAILS

Marketplace-specific prompts	Context injection
No SQL generation	No guessing
Explicit “data not available”	Scope / forbidden-topic filters
Quota management	Fallback responses

HUMAN APPROVAL BOUNDARY



SECTION 05 / VALIDATION

05

Shipping the Experience


From roadmap to production: what shipped, why it shipped, and how users experience it.

Roadmap

Production UI

Live Demo

Roadmap: build trust before intelligence




✓ SHIPPED · ○ PLANNED Every phase unlocked the next - we only introduced AI **after** users trusted the underlying profit engine.

AI Business Insights Chat

Live production screen

The AI layer became the fastest path from “**what happened?**” to “**what should I do next?**” – replacing hours of spreadsheet analysis with grounded, conversational recommendations.



- 1 **Quota indicator** – remaining AI insights, always visible.
- 2 **Voice toggle** – hands-free queries & spoken replies.
- 3 **Grounded starters** – seeded from settlement-aligned data.
- 4 **Voice input** – type or tap the mic (en-IN).

WHY IT'S NOT "CHATGPT ON DATA"

Every response is grounded in **pre-computed marketplace analytics** – not live SQL or vector retrieval.


BUSINESS QUESTIONS ANSWERED

Why did profit drop? · Which SKU to discontinue? · Where is money leaking? · Which returns need attention?

15	100	250
Trial · insights/mo	Basic · insights/mo	Pro · insights/mo

True profit, at a glance

Live production screen




- 1 True Profit is the hero**
The dark card anchors the metric row – the number no report gives directly, after every fee and COGS.
- 2 Actionable insights before analytics**
Quick Insights surface settlement ratio, fee burden and loss-making SKUs above the fold.
- 3 Every metric is explorable**
One click from summary to SKU, geography, returns and advertising detail.
- 4 AI is grounded in this same view**
Chat answers reconcile with the dashboard because both share one business context.

DASHBOARD PRINCIPLES

- Profit first
- Insights second
- Charts third
- Settlement-aligned • no vanity metrics

From seller problem to shipped product



THE THROUGH-LINE Every stage earned the next — we only shipped AI once the underlying profit engine was trusted.

Beyond AI Chat: the complete SellerOps workflow

Each module solves a different seller problem while feeding the same profit engine.

← Back to tools



Welcome to Seller Analytics!

Let's set up your account by importing your Amazon Seller data. This will only take a few minutes.

What you'll need:

1 Orders

Amazon All Orders Report

2 Transactions

Payment report from Reports Repository

3 COGS

Cost of Goods Sold

DI Skip for now

Get Started →

STEP 01 • IMPORT

Onboarding

Import Amazon & Meesho reports in minutes.

Faster setup

Amazon Analytics

Analytics Dashboard - Pro

Dashboard

AI Analysis

Product Performance

Orders

COGS

Inventory

Analytics

Reconciliation

Reports

Settings

Karandeep Singh

Account Settings

← Back to tools

Cost of Goods Sold (COGS)

62 SKUs sold in April 2026 • Enter product costs to calculate profit

New SKUs start at 0 cost. Download the template, fill your costs, and upload it.

17 SKUs still need a product cost. Please enter all product costs to generate accurate profit reports.

Search Products

PRODUCT NAME

Pioneer Trek Set Mini 142

SKU: YRV-8733 ASIN: B015YZ7ZB8

PRODUCT NAME

Atlas Utility Kit Go 148

SKU: FRG-5458 ASIN: B081B625V2

PRODUCT NAME

Summit Comfort Set Trio 109

SKU: LUM-5081 ASIN: B0PH8R7TH4

PRODUCT NAME

Cobalt Home Essential Duo 111

SKU: HBR-3065 ASIN: B0WN6NNWSQ

PRODUCT NAME

Quill Active Bundle Solo 141

SKU: WLW-4169 ASIN: B0E4JN9SSG

PRODUCT NAME

Nova Utility Kit Duo 146

SKU: CAS-3566 ASIN: B08WAPQ1MQ

STEP 02 • CONFIGURE

COGS management

Ensure every profit calculation uses real product costs.

Trusted calculations

Amazon Analytics

Analytics Dashboard - Pro

Dashboard

AI Analysis

Product Performance

Orders

COGS

Inventory

Analytics

Profit Analytics

Ads Analytics

Returns Analytics

Geo Analytics

Reconciliation

Reports

Settings

Karandeep Singh

Account Settings

← Back to tools

Profit & Loss Summary

580 transactions analyzed

Gross Revenue

₹2,09,190.19

Refunds

-₹14,329.06

Net Revenue

₹1,94,861.13

Total Fees

-₹98,600.1

COGS (Actual)

-₹50,123.38

ADA + Ad Spend

-₹35,667.99

Net Profit

₹49,207.7

Detailed P&L Statement

Click to collapse

REVENUE

Product Sales

Shipping Credits

Gross Revenue

DEDUCTIONS & RETURNS

Refunds (38 orders)

Net Revenue

STEP 03 • ANALYZE

True Profit Dashboard

See true profit after every fee, refund and ad cost.

Better decisions

Amazon Analytics

Analytics Dashboard - Pro

Dashboard

AI Analysis

Product Performance

Orders

COGS

Inventory

Analytics

Profit Analytics

Ads Analytics

Returns Analytics

Geo Analytics

Reconciliation

Reports

Settings

Karandeep Singh

Account Settings

← Back to tools

Order Reconciliation

Track matching between orders and settlement transactions

Total Orders

601

Matched

601

100% match rate

Unmatched

0

Reconciliation Progress

601 matched

100% Matched

Overview

Matched (601)

Unmatched (0)

Reconciliation Flow

Orders Imported

601 total orders

Transactions Matched

Linked by Order ID

Settlement Calculated

601 orders reconciled

Status Summary

Switch

46

Unmatched

0

STEP 04 • VERIFY

Reconciliation

Explain exactly where every rupee went.

Verified settlements

Every module feeds the same trusted profit engine that powers SellerOps AI

05 • Build & Experience

WORKING DEMO & WALKTHROUGH

Don't take my word for it - try it




INTERACTIVE DEMO · STORYLANE

Scan to click through the live product

No signup, no data upload — a guided walkthrough of the real dashboard and AI chat.

demo.storylane.com/share/7zvkbub3sji

LOOM WALKTHROUGH



Scan for the narrated tour

loom.com/share/43dfbb83...

SECTION 06 / IMPACT & LEARNINGS

06

Outcomes & Learnings

The product shipped successfully. Here's what we've validated, what remains unmeasured, and what I'd do next.

What shipped

What I learned

Future roadmap

Early validation, long-term measurement pending

WHAT IS KNOWN

- ✓ AI Business Insights shipped to production
- ✓ Used across 30+ seller accounts
- ✓ Processes 2L+ marketplace orders
- ✓ Built on ₹60L+ GMV of business data

30+

Seller accounts connected

2L+

Marketplace orders processed

₹60L+

GMV analysed

NEXT PHASE · OUTCOME MEASUREMENT

ADOPTION

- AI activation rate
- Weekly active AI users
- Questions per user

QUALITY

- Dashboard match %
- Hallucination rate
- Rec. acceptance

BUSINESS

- Upgrade conversion
- Retention
- Plan expansion

Next: measure activation, engagement & business impact before expanding AI capabilities.

BIGGEST TAKEAWAY

The biggest risk was never building the AI - it was ensuring **every answer matched the seller’s financial dashboard**. That trust layer shipped first; product analytics come next.

Business impact: why this product matters

WITHOUT SELLEROPS

- Manual reconciliation across multiple reports
- True profit hidden by fees, returns & ads
- Loss-making SKUs run unnoticed for weeks
- Decisions delayed by spreadsheet work

BEFORE 5 different reports → AFTER 1 dashboard + AI chat

VALUE CREATED

- ✓ Understand profit without Excel
- ✓ Explain *why* profit changed — not just that it did
- ✓ Discover loss-making SKUs earlier
- ✓ Decide in minutes instead of hours
- ✓ AI answers always reconcile with the dashboard

Next: instrument adoption & business impact now that the product is live.

OUTCOME **SellerOps shortens the distance between raw marketplace data and confident business decisions.**

Five things this project taught me

- 01 In seller tools, trust is the product – and trust is math.** Every incorrect number destroys trust faster than any missing feature.
- 02 Onboarding is part of the product.** The hardest step isn't signing up – it's getting trustworthy data in.
- 03 AI differentiation only works when grounded.** The advantage wasn't the model – it was the context, guardrails and business logic.
- 04 One product can serve multiple markets with different workflows.** Amazon & Meesho needed different onboarding, prompts and analytics on one profit engine.
- 05 Subscription UX is a retention decision.** Pricing and quotas shape behavior as much as features do – monetization is UX.

BIGGEST TAKEAWAY Reliable business software demanded far more product thinking around trust, onboarding and workflows than around AI models themselves.

What's next - and why

P0	Activation funnel optimization	Upload friction limits activation more than feature depth.
P0	Instrument validation metrics	Instrument activation, engagement & recommendation quality before scaling AI.
P1	Golden evaluation dataset	Build offline regression tests before expanding autonomous AI capabilities.
P1	Data freshness automation	Notify sellers when reports go stale — preserves AI quality, drives re-engagement.
P1	Scale the Meesho growth engine	The active acquisition bet - ship Meesho onboarding ahead of Amazon parity.
P2	Improve upgrade moments	Surface premium insights at the right moment to increase paid conversion.
P3	Listing Generator: validate or remove	Run a fake-door experiment. If demand stays low, remove it.

THE PRINCIPLE Every item removes a constraint discovered while shipping — the roadmap follows evidence, not feature requests.

Evidence of execution

PRODUCT

- ✓ Live SellerOps platform
- ✓ Active seller onboarding
- ✓ Production AI insights
- ✓ Interactive product demo

DESIGN

- ✓ Complete Figma system
- ✓ User flows & wireframes
- ✓ Design iterations
- ✓ Production screenshots

ENGINEERING

- ✓ Production architecture
- ✓ AI pipeline & guardrails
- ✓ Working AI workflow
- ✓ Grounded context layer

VALIDATION

- ✓ Early seller usage signal
- Product analytics rollout
- Loom walkthrough
- Golden eval dataset

SellerOps isn't a concept — it's a production product I researched, designed, built, shipped, and continue to improve.



SellerOps AI

FINAL THOUGHT

Building successful AI products isn't about smarter models. It's about products that **users trust.**

SellerOps taught me that building AI products is less about choosing models and more about understanding users, earning trust, and making difficult product trade-offs. This project reflects that end-to-end journey - from discovering the problem to shipping a production system.

